

Connectionist Computing

COMP 30230/41390

Gianluca Pollastri

office: E0.95, Science East.

email: gianluca.pollastri@ucd.ie

Credits

- **Geoffrey Hinton, University of Toronto.**
 - borrowed some of his slides for “Neural Networks” and “Computation in Neural Networks” courses.



- **Ronan Reilly, NUI Maynooth.**
 - slides from his CS4018.



- **Paolo Frasconi, University of Florence.**
 - slides from tutorial on Machine Learning for structured domains.



Lecture notes on Brightspace

- **Strictly confidential...**
- **Slim PDF version will be uploaded later, typically the same day as the lecture.**
- **If there is demand, I can upload onto Brightspace last year's narrated slides.. (should be very similar to this year's material)**

Books

- No book covers large fractions of this course.
- Parts of chapters 4, 6, (7), 13 of Tom Mitchell's "Machine Learning"
- Parts of chapter V of Mackay's "Information Theory, Inference, and Learning Algorithms", available online at:
<http://www.inference.phy.cam.ac.uk/mackay/itprnn/book.html>
- Chapter 20 of Russell and Norvig's "Artificial Intelligence: A Modern Approach", also available at:
<http://aima.cs.berkeley.edu/newchap20.pdf>
- More materials later..

Marking

- 3 landmark papers to read, and submit a 10-line summary on Brightspace about: each worth 6-7%
- a connectionist model to build and play with on some sets, write a report: 30%
- Final Exam in the RDS (50%)

Assignment 1

- Read the first section of the following article by Marvin Minsky:
- <http://web.media.mit.edu/~minsky/papers/SymbolicVs.Connectionist.html>
- down to “.. we need more research on how to combine both types of ideas.”
- Submit through BrightSpace a 250 word MAX summary by October 2nd at 23:59, any time zone of your choice (Baker Island?).
- One third of a grade down every 2 days late, that is: if you deserve an A and you're 1-2 days late you get an A-, 3-4 days late a B+, etc.
- Make sure it's gone through!


Gradient descent and associators

- Complexity of gradient computation is minimal: $O(nm)$
- BUT, it is unclear how many steps along the gradient need to be taken...
- Much better storage of examples than with one-shot learning, though typically there will be residual error.

$$\Delta w_{ji} = -\eta \sum_{p=1}^P (y_j^{(p)'} - y_j^{(p)}) x_i^{(p)} =$$
$$\eta \sum_{p=1}^P (y_j^{(p)} - y_j^{(p)'}) x_i^{(p)}$$

Hopfield Nets

- Networks of binary threshold units.
- Feedback networks: each units has connections to all other units except itself.



The energy function

- The global energy is the sum of many contributions. Each contribution depends on one connection weight and the binary states of **two** neurons:

$$E = -\frac{1}{2} \sum_{i,j} w_{ij} y_i y_j - \frac{1}{2} \sum_i b_i y_i$$

- The simple energy function makes it easy to compute how the state of one neuron affects the global energy (it is the activation of neuron!):

$$E(y_i = -1) - E(y_i = 1) = \sum_j w_{ij} y_j + b_i$$

Storing memories (learning)

- If we want to store a set of memories:

$$y^{(1)}, \dots, y^{(p)}, \dots, y^{(P)}$$
$$y^{(p)} = (y_1^{(p)}, \dots, y_m^{(p)})$$


- if the states are -1 and $+1$ then we can use the update rule:

$$\Delta w_{ji} = \eta \sum_p y_i^{(p)} y_j^{(p)}$$

Example

- Say we have 4 patterns of size 4:

- (1, -1, -1, -1)
- (-1, 1, -1, -1)
- (-1, -1, 1, -1)
- (-1, -1, -1, 1)



- We build the sigmoid Hopfield net based on the 4 patterns (Matlab nnet toolbox).
- Incidentally, here one-shot learning would not work: try.

$$\Delta w_{ji} = \eta \sum_p y_i^{(p)} y_j^{(p)}$$

Example

- Let's now start from some state Y for the neurons, and watch the network evolve.
- $Y=(1 \ 0 \ 0 \ 0)$

Steps (1 update/neuron):

1: (0.4999 -0.6620 -0.6620 -0.6620)

2: (0.5022 -0.8476 -0.8476 -0.8476)

3: (0.6351 -0.9332 -0.9332 -0.9332)

4: (0.8186 -1.0000 -1.0000 -1.0000)

5: (1 -1 -1 -1) converged to pattern 1!



Example (2)

- Different starting point:
- $Y=(0\ 0\ 0\ 0)$

Steps:

1: (-0.4273 -0.4273 -0.4273 -0.4273)

2: (-0.5226 -0.5226 -0.5226 -0.5226)

3: (-0.5439 -0.5439 -0.5439 -0.5439)

4: (-0.5486 -0.5486 -0.5486 -0.5486)


..

7: (-0.55 -0.55 -0.55 -0.55)

stuck in the middle!



Hopfield Network State Space



Example (3)

- Let's now try train a Hopfield net on slightly nastier vectors (the previous ones were orthogonal):

(1.0000 -1.0000 -1.0000 0.3000)
(0.1000 1.0000 -1.0000 -0.1000)
(-1.0000 -1.0000 1.0000 -0.5000)
(-1.0000 -1.0000 -1.0000 1.0000)

Example (3)

- Let's start from some state Y for the neurons, and watch the network evolve.
- $Y=(1\ 0\ 0\ 0)$

Steps (1 update/neuron):

1:	(0.9957	-0.3693	-0.3693	-0.0028)	
5:	(1.0000	-0.5332	-0.5332	-0.0040)	
50:	(1.0000	-0.5348	-0.5348	-0.0040)	odd plateau
170:	(1.0000	-0.5345	-0.5350	-0.0040)	
5000:	(1.0000	0.2032	-1.0000	1.0000)	spurious min!

From Hopfield nets to Boltzmann machine

- Hopfield nets (attempt to) minimise an energy function:

$$E(y) = -\frac{1}{2} \sum_{i,j} w_{ij} y_i y_j = -\frac{1}{2} \mathbf{y}^T \mathbf{W} \mathbf{y}$$

$$\mathbf{y} = (y_1, \dots, y_m)^T$$

$$\mathbf{W} = \begin{pmatrix} w_{11} & \dots & w_{1i} & \dots & w_{1n} \\ w_{j1} & \dots & w_{ji} & \dots & w_{jn} \\ w_{m1} & \dots & w_{mi} & \dots & w_{mn} \end{pmatrix}$$

Stochastic units

- Replace the binary threshold units by binary *stochastic* units. A neuron is switched on/off with a certain probability instead of deterministically.

Stochastic units


- **Reminder: the energy reduction by switching a neuron on is:**

$$z_i = \sum_j w_{ij} y_j = E(y_i = -1) - E(y_i = 1)$$

Stochastic units

- The probability that a given unit is switched on is a function of the amount of energy that its change of state would contribute to the network's overall energy, and the "temperature" T .

$$\frac{\exp(z_i/T)}{\exp(z_i/T) + \exp(-z_i/T)} = \frac{1}{1 + \exp(-2z_i/T)}$$




Hopfield vs. Boltzmann

- Since in Boltzmann networks the unit update rule (or activation function) has a probabilistic component, if we allow the network to run it will settle into a given equilibrium state only with a certain probability.
- Hopfield networks, on the other hand, are deterministic. Given an initial state their equilibrium state is determined.

Stochastic units

- **Temperature makes it easier to cross energy barriers.**
 - **Start at high temperature where its easy to cross energy barriers.**
 - **Reduce slowly to low temperature where good states are much more probable than bad ones.**



Implementing a probability distribution

- It can be shown that the Boltzmann machine generates configurations according to the distribution:

$$P(\mathbf{y}|\mathbf{W}) = \frac{1}{Z(\mathbf{W})} \exp\left(\frac{1}{2}\mathbf{y}^T \mathbf{W} \mathbf{y}\right)$$

- Where:

$$Z(\mathbf{W}) = \sum_{\mathbf{y}} \exp\left(\frac{1}{2}\mathbf{y}^T \mathbf{W} \mathbf{y}\right)$$

Learning in the Boltzmann machine

- Working on the probability distribution it is possible to devise the following learning rule:

$$\Delta w_{ij} = \eta \sum_{p=1}^P \left[y_i^{(p)} y_j^{(p)} - \sum_{\mathbf{y}} y_i y_j P(\mathbf{y} | \mathbf{W}) \right]$$

Learning in the Boltzmann machine

- **First term:**

$$\eta \sum_{p=1}^P y_i^{(p)} y_j^{(p)}$$

- **Empirical correlation between neurons i and j, measured from the examples (same as learning rule for Hopfield).**

Learning in the Boltzmann machine

- **Second term:**

$$\eta \sum_{\mathbf{y}} y_i y_j P(\mathbf{y} | \mathbf{W})$$

- **Correlation between neurons i and j, measured on patterns generated according to the probability distribution underlying the model. Notice that we're summing over all possible patterns (2^N)**

Learning in the Boltzmann machine

- **The first term is readily evaluated from the examples.**
- **The second term can be estimated by letting the model evolve until equilibrium, measuring the correlations, and repeating many times: can be computationally tough.**
- **Monte Carlo.**

Interpretation of learning terms

- **First term: the network is awake and measures the correlations in the real world.**
- **Second term: the network sleeps and dreams about the world using the model it has of it.**
- **Once dream and reality coincide learning reaches an end. It is interesting to notice that the network *unlearns* its dreams.**

Weakness of Hopfield nets and Boltzmann machines

- All units are *visible*, i.e. correspond to observable stuff (components of the examples).
- In this situation nothing more than second order interactions can be captured by Hopfield nets and the Boltzmann machine.
- If for instance the examples are bits of images, second order statistics are a poor representation.

Hidden units for Hopfield nets/Boltzmann machines

- Instead of using the net just to store memories, use it to construct *interpretations* of the input.
 - The input is represented by the visible units.
 - The interpretation is represented by the states of the hidden units.
- Higher order correlations can be represented by hidden units.
- More powerful model, but even harder to train.

