

Connectionist Computing

COMP 30230/41390

Gianluca Pollastri

office: E0.95, Science East.

email: gianluca.pollastri@ucd.ie

Credits

- **Geoffrey Hinton, University of Toronto.**
 - borrowed some of his slides for “Neural Networks” and “Computation in Neural Networks” courses.



- **Ronan Reilly, NUI Maynooth.**
 - slides from his CS4018.



- **Paolo Frasconi, University of Florence.**
 - slides from tutorial on Machine Learning for structured domains.



Lecture notes on Brightspace

- **Strictly confidential...**
- **Slim PDF version will be uploaded later, typically the same day as the lecture.**
- **If there is demand, I can upload onto Brightspace last year's narrated slides.. (should be very similar to this year's material)**

Books

- No book covers large fractions of this course.
- Parts of chapters 4, 6, (7), 13 of Tom Mitchell's "Machine Learning"
- Parts of chapter V of Mackay's "Information Theory, Inference, and Learning Algorithms", available online at:
<http://www.inference.phy.cam.ac.uk/mackay/itprnn/book.html>
- Chapter 20 of Russell and Norvig's "Artificial Intelligence: A Modern Approach", also available at:
<http://aima.cs.berkeley.edu/newchap20.pdf>
- More materials later..

Marking

- 3 landmark papers to read, and submit a 10-line summary on Brightspace about: each worth 6-7%
- a connectionist model to build and play with on some sets, write a report: 30%
- Final Exam in the RDS (50%)

Learning in the Boltzmann machine with hidden units


- Even in this case we have two terms in Δw_{ij} .
- The first term is the correlation between neurons i and j when the visible units are clamped to the examples.
- The second is the correlation between neurons i and j when the system is let evolve freely.

Learning in the Boltzmann machine with hidden units

- **In this case both terms must be estimated by letting the network evolve many times until equilibrium, in one case with the visible units clamped, in the other freely.**
- **Not only: the missing units need to be estimated separately for each example.**
- **Can be computationally very expensive**

Restricted Boltzmann Machines (RBM)

- We restrict the connectivity to make inference and learning easier.
 - Only one layer of hidden units.
 - No connections between hidden units.
- It only takes one step to reach equilibrium when the visible units are clamped..



Types of learning task

- **Supervised learning:**
 - learning to predict output when examples of input and corresponding output are available.
- **Reinforcement learning:**
 - no real supervision, only occasional payoff: e.g. I don't know if a chess move is correct but I know if I win.
- **Unsupervised learning:**
 - I want to create an internal representation of the data, e.g. in form of clusters, extract relevant features for further tasks, etc.

Supervised learning

- **A set of examples: $\langle x, f(x) \rangle$**
 - **x is some object (instance) $\in \mathcal{X}$ (instance space)**
 - **$f(\cdot)$ is an *unknown* function**
- **Learning algorithm will guess $h(\cdot) \approx f(\cdot)$**
- **Inductive learning hypothesis**
 - **Any $h(\cdot)$ that approximates $f(\cdot)$ well on training examples will also approximate $f(\cdot)$ well on new (unseen) instances x . This isn't necessarily true, of course..**

Learning as refinement

- **Start with a small hypothesis class (e.g., boolean conjunctions)**
 - this means we need to *know a priori* something about the solution
- **Use examples to infer the particular function (e.g. conjunction) in the class.**


Learning as refinement

- **Start with a small hypothesis class (e.g., boolean conjunctions)**
 - this means we need to *know a priori* something about the solution
- **Use examples to infer the particular function (e.g. conjunction) in the class.**

Generalisation

- So far we discussed how well we learn the training set.
- What really matters is how well we *generalise* i.e. how well we perform on instances not seen before.

Generalisation



$$\text{true error} = \int_{\mathcal{X}} f(x) \oplus h(x) P_D(x) dx$$

Generalisation

- **Are there boundaries on generalisation error?**
- **How expressive is a model, i.e. once we know how to learn:**
 - **what can we learn?**
 - **how many examples do we need to learn?**
- **For instance: how complex a task can we learn with N neurons / how many examples do we need to train them?**

PAC learning

- **Probably Approximately Correct (hopefully)**
- **H finite, unknown data distribution D, consistent learner**
- **Then we need at least**

$$m \geq \frac{1}{\varepsilon} (\ln |H| + \ln(1/\delta))$$

training examples to guarantee that the true error will be $< \varepsilon$ with probability $> (1-\delta)$

m is called *sample complexity*

Example

- 1000 functions in H , $\delta=1\%$, $\epsilon=1\%$
- $m \geq 100 (\ln 1000 + \ln 100) = 1151.3$

Examples

Hypothesis space 1: $|H|=2^{2^n}$

$$m \geq \frac{1}{\varepsilon} (2^n \ln 2 + \ln(1/\delta))$$

Intractable!!


Hypothesis space 2: $|H|=3^n$

$$m \geq \frac{1}{\varepsilon} (n \ln 3 + \ln(1/\delta))$$

Polynomial


Sample complexity bound

$$m \geq \frac{1}{\varepsilon} (n \ln 3 + \ln(1/\delta))$$



Binary threshold neuron

$$z = \sum_i x_i w_i$$
$$y = \begin{cases} 1 & \text{if } z > \theta \\ 0 & \text{otherwise} \end{cases}$$




- **How powerful is it, how many examples do we need? Can we apply PAC boundaries?**


Perceptron

- PAC works well if we have a limited space H . But for a perceptron the hypothesis space is the set of all separation hyperplanes
- How many hyperplanes in \mathbb{R}^n ? $|H|=\infty \dots ?$
- Not only, but PAC is based on consistent learners. We know that perceptrons can't learn some tasks.
- We need an extension of PAC:
- VC dimension
V=Vapnik
C=Chervonenkis

Shattering

- H shatters a set of points S if for every possible dichotomy of S (way of splitting it in 2), there exists a consistent hypothesis h in H
- Example: is this set of points shattered by the hypothesis space of all axis parallel rectangles?






VC dimension

- **VC(H) is the size of the largest set of points S that can be shattered by H**
- **Examples for 2D data points:**
 - **VC(axis-parallel-rectangles)=4**
 - **VC(circles)=3**
 - **VC(lines)=3**


VC dimension

- For n -dimensional data points a linear decision surface has $VC(H) = n+1$
- For instance, for a perceptron with n inputs (n -dimensional data points), $VC(H)=n+1$



PAC learning and VC dim

$$m \geq \frac{1}{\varepsilon} (4 \log_2(2/\delta) + 8VC(H) \log_2(13/\varepsilon))$$



PAC learning for a perceptron

- n inputs $\rightarrow VC(H)=n+1$
- Example:
- 100 inputs, $\delta=1\%$, $\epsilon=1\%$
- $m \geq 100 (4 \log 200 + 8 \cdot 10^1 \log 1300) = 8388.8$

- **Enough about learning theory for the moment.**
- **We will get back to it soon to compute the expressive power of a new kind of network..**